

# CS6023: GPU Programming

Assignment-2 (15 marks)

Deadline: March 5th, 2023, 23:55 on Moodle

## 1 Problem Statement

Given four integer matrices  $A_{p \times q}$ ,  $B_{q \times r}$ ,  $C_{p \times q}$ , and  $D_{r \times q}$ , efficiently compute matrix  $E$

$$E = AB + CD^T \tag{1}$$

by considering the aspects of **memory coalescing** and **shared memory**.

## 2 Input and Output

### 2.1 Input format

- First line of input contain 3 integers  $p$ ,  $q$ , and  $r$ .
- Next  $p$  lines contain  $q$  space-separated integers representing matrix  $A$ .
- Next  $q$  lines contain  $r$  space-separated integers, representing matrix  $B$ .
- Next  $p$  lines contain  $q$  space-separated integers representing matrix  $C$ .
- Next  $r$  lines contain  $q$  space-separated integers representing matrix  $D$ .

### 2.2 Output format

- Write  $p$  lines each containing  $r$  space separated integers representing matrix  $E_{p \times r}$

### 2.3 Constraints

- $2 \leq p, q, r \leq 1024$ .
- Each integer of matrices  $A, B, C, D$  is in range  $[-10, 10]$ .

### 3 Sample Testcase

- Let  $p = 2, q = 3, r = 2$

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} B = \begin{bmatrix} 1 & -2 \\ -4 & 3 \\ 5 & -6 \end{bmatrix} C = \begin{bmatrix} 4 & 2 & 10 \\ 7 & 5 & -8 \end{bmatrix} D = \begin{bmatrix} 8 & 1 & 5 \\ -3 & 3 & -6 \end{bmatrix}$$

$$\implies AB = \begin{bmatrix} 8 & -14 \\ 14 & -29 \end{bmatrix} \quad CD^T = \begin{bmatrix} 84 & -66 \\ 21 & 42 \end{bmatrix}$$

$$\implies E = AB + CD^T = \begin{bmatrix} 92 & -80 \\ 35 & 13 \end{bmatrix}$$

### 4 Points to be noted

- The file `main.cu` provided by us contains the code, which takes care of taking the input, printing the result, and printing the execution time.
- Don't write any code in the `main()` function. Do not write any print statements.
- You need to implement the `compute()` function provided in the `main.cu`.
- You are free to use any number of functions/kernels.
- You can launch the kernels as you wish.
- Test your code on large input matrices.
- **It is compulsory to optimize for coalesced accesses. Also, make use of shared memory.**

### 5 Submission Guidelines

- Rename `main.cu` to `YourRollNumber.cu`. For example, if your roll number is `cs19b100`, then rename the file to `cs19b100.cu`.
- **Submit only this .cu file on Moodle.**
- After submission, download the file and make sure it was the one you intended to submit.

## 6 Learning suggestions

- Write a CPU version of code achieving the same functionality. Time the CPU code and GPU code separately for large matrices and compare the performances.
- Try to use a minimum number of kernel calls.
- Try to use `syncthreads()` and `syncwarps()` as minimum as possible to gain performance benefits.
- Exploit shared memory as much as possible to gain performance benefits.
- Think of an approach or technique that can avoid bank conflicts of shared memory.
- To understand how good coalescing and shared-memory utilization are, use Nvidia profilers to see effective memory bandwidth on large test cases.